

BUISNESS REPORT – Linear Regression:

The data is taken from the comp-activ databases which is a collection of computer systems activity measures. The data was collected from a Sun Sparcstation 20/712 with 128 Mbytes of memory running in a multi-user university department. The data was collected continuously every 5 seconds. Using the various system programs which are running in the background for every task being performed by the user, Predict the percentage portion of time (out of 100), that cpu runs in user mode, and how does each system program affect the same.

Tasks to be performed:

1. Load data and describe data
 - a. Import necessary libraries & packages
 - b. Load dataset
 - c. Check necessary details about data like shape, data types of the variable, missing values etc.

```
import pandas as pd
import numpy as np
import seaborn as sns
import seaborn as sb
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from statsmodels.stats.outliers_influence import variance_inflation_factor
import math
from sklearn.preprocessing import MinMaxScaler
import warnings
warnings.filterwarnings("ignore")
```

Loading the dataset

```
df = pd.read_csv('compactiv.csv')
df.head()
```

From the table we can see all the required variables that we require excel sheet

	0	1	2	3	4
lread	1	0	15	0	5
lwrite	0	0	3	0	1
scall	2147	170	2162	160	330
sread	79	18	159	12	39
swrite	68	21	119	16	38
fork	0.2	0.2	2.0	0.2	0.4
exec	0.2	0.2	2.4	0.2	0.4
rchar	40671.0	448.0	NaN	NaN	NaN
wchar	53995.0	8385.0	31950.0	8670.0	12185.0
pgout	0.0	0.0	0.0	0.0	0.0
ppgout	0.0	0.0	0.0	0.0	0.0
pgfree	0.0	0.0	0.0	0.0	0.0
pgscan	0.0	0.0	0.0	0.0	0.0
atch	0.0	0.0	1.2	0.0	0.0
pgin	1.6	0.0	6.0	0.2	1.0
ppgin	2.6	0.0	9.4	0.2	1.2
pflt	16.0	15.63	150.2	15.6	37.8
vflt	26.4	16.83	220.2	16.8	47.6
runqsz	CPU_Bound	Not_CPU_Bound	Not_CPU_Bound	Not_CPU_Bound	Not_CPU_Bound
freemem	4670	7278	702	7248	633
freeswap	1730946	1869002	1021237	1863704	1760253
usr	95	97	87	98	90

Check the necessary detail about data is given by its info, shape, data types:

```
df.shape
```

```
(8192, 22)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 8192 entries, 0 to 8191
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	lread	8192 non-null	float64
1	lwrite	8192 non-null	float64
2	scall	8192 non-null	float64
3	sread	8192 non-null	float64
4	swrite	8192 non-null	float64
5	fork	8192 non-null	float64
6	exec	8192 non-null	float64
7	rchar	8192 non-null	float64
8	wchar	8192 non-null	float64
9	pgout	8192 non-null	float64
10	ppgout	8192 non-null	float64
11	pgfree	8192 non-null	float64
12	pgscan	8192 non-null	float64
13	atch	8192 non-null	float64
14	pgin	8192 non-null	float64
15	ppgin	8192 non-null	float64
16	pflt	8192 non-null	float64
17	vflt	8192 non-null	float64
18	freemem	8192 non-null	float64
19	freeswap	8192 non-null	float64
20	usr	8192 non-null	float64
21	CPU_Bound_CPU_Bound	8192 non-null	uint8
22	CPU_Bound_Not_CPU_Bound	8192 non-null	uint8

```
dtypes: float64(21), uint8(2)
```

```
memory usage: 1.4 MB
```

df.dtypes

```
lread      int64
lwrite     int64
scall      int64
sread      int64
swrite     int64
fork       float64
exec       float64
rchar      float64
wchar      float64
pgout      float64
ppgout     float64
pgfree     float64
pgscan     float64
atch       float64
pgin       float64
ppgin      float64
pflt       float64
vflt       float64
runqsz     object
freemem    int64
freeswap   int64
usr        int64
dtype: object
```

```
: df.isnull().sum()
```

```
: lread      0
   lwrite     0
   scall      0
   sread      0
   swrite     0
   fork       0
   exec       0
   rchar     104
   wchar      15
   pgout      0
   ppgout     0
   pgfree     0
   pgscan     0
   atch       0
   pgin       0
   ppgin      0
   pflt       0
   vflt       0
   runqsz     0
   freemem    0
   freeswap   0
   usr        0
dtype: int64
```

From the above code and result we can figure out if there any null values, missing values, datatype of all the variables. Rchar has 104 missing values and wchar has 15 missing values.

2. Perform EDA and data cleaning

a. Generate the summary statistics for each of the variables and write comments on your observations

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
lread	8192.0	1.955969e+01	53.353799	0.0	2.0	7.0	20.000	1845.00
lwrite	8192.0	1.310620e+01	29.891726	0.0	0.0	1.0	10.000	575.00
scall	8192.0	2.306318e+03	1633.617322	109.0	1012.0	2051.5	3317.250	12493.00
sread	8192.0	2.104800e+02	198.980146	6.0	86.0	166.0	279.000	5318.00
swrite	8192.0	1.500582e+02	160.478980	7.0	63.0	117.0	185.000	5456.00
fork	8192.0	1.884554e+00	2.479493	0.0	0.4	0.8	2.200	20.12
exec	8192.0	2.791998e+00	5.212456	0.0	0.2	1.2	2.800	59.56
rchar	8088.0	1.973857e+05	239837.493526	278.0	34091.5	125473.5	267828.750	2526649.00
wchar	8177.0	9.590299e+04	140841.707911	1498.0	22916.0	46619.0	106101.000	1801623.00
pgout	8192.0	2.285317e+00	5.307038	0.0	0.0	0.0	2.400	81.44
ppgout	8192.0	5.977229e+00	15.214590	0.0	0.0	0.0	4.200	184.20
pgfree	8192.0	1.191971e+01	32.363520	0.0	0.0	0.0	5.000	523.00
pgscan	8192.0	2.152685e+01	71.141340	0.0	0.0	0.0	0.000	1237.00
atch	8192.0	1.127505e+00	5.708347	0.0	0.0	0.0	0.600	211.58
pgin	8192.0	8.277960e+00	13.874978	0.0	0.6	2.8	9.765	141.20
ppgin	8192.0	1.238859e+01	22.281318	0.0	0.6	3.8	13.800	292.61
pfit	8192.0	1.097938e+02	114.419221	0.0	25.0	63.8	159.600	899.80
vfit	8192.0	1.853158e+02	191.000603	0.2	45.4	120.4	251.800	1365.00
freemem	8192.0	1.763456e+03	2482.104511	55.0	231.0	579.0	2002.250	12027.00
freeswap	8192.0	1.328126e+06	422019.426957	2.0	1042623.5	1289289.5	1730379.500	2243187.00
usr	8192.0	8.396887e+01	18.401905	0.0	81.0	89.0	94.000	99.00

We can see from the above summary statistics the mean of - Number of characters transferred per second by system write calls is the maximum. Most of the variables have minimum values from 0. The maximum Number of characters transferred per second by system read calls is 2526649 which is the highest frequency reached among all other variables.

b. Working with Null/Missing values

i. Check for Missing values and perform the necessary steps for data imputation and provide reasoning for your approach.

As we saw in the previous question there are missing values that must be replaced. Since both are numerical in nature we must replace them by the median.

```
df['rchar'].fillna(df.rchar.median(), inplace = True)
df
```

	lread	lwrite	scall	sread	swrite	fork	exec	rchar	wchar	pgout	...	pgscan	atch	pgin	ppgin	pfit	vfit	runqsz	freemem	freesv
0	1	0	2147	79	68	0.2	0.20	40671.0	53995.0	0.00	...	0.00	0.0	1.60	2.60	16.00	26.40	CPU_Bound	4670	1730
1	0	0	170	18	21	0.2	0.20	448.0	8385.0	0.00	...	0.00	0.0	0.00	0.00	15.63	16.83	Not_CPU_Bound	7278	1869
2	15	3	2162	159	119	2.0	2.40	125473.5	31950.0	0.00	...	0.00	1.2	6.00	9.40	150.20	220.20	Not_CPU_Bound	702	1021
3	0	0	160	12	16	0.2	0.20	125473.5	8670.0	0.00	...	0.00	0.0	0.20	0.20	15.60	16.80	Not_CPU_Bound	7248	1863
4	5	1	330	39	38	0.4	0.40	125473.5	12185.0	0.00	...	0.00	0.0	1.00	1.20	37.80	47.60	Not_CPU_Bound	633	1760
...
8187	16	12	3009	360	244	1.6	5.81	405250.0	85282.0	8.02	...	55.11	0.6	35.87	47.90	139.28	270.74	CPU_Bound	387	986
8188	4	0	1596	170	146	2.4	1.80	89489.0	41764.0	3.80	...	0.20	0.8	3.80	4.40	122.40	212.60	Not_CPU_Bound	263	1055
8189	16	5	3116	289	190	0.6	0.60	325948.0	52640.0	0.40	...	0.00	0.4	28.40	45.20	60.20	219.80	Not_CPU_Bound	400	969
8190	32	45	5180	254	179	1.2	1.20	62571.0	29505.0	1.40	...	18.04	0.4	23.05	24.25	93.19	202.81	CPU_Bound	141	1022
8191	2	0	985	55	46	1.6	4.80	111111.0	22256.0	0.00	...	0.00	0.2	3.40	6.20	91.80	110.00	CPU_Bound	659	1756

8192 rows x 22 columns

```
df['wchar'].fillna(df.wchar.median(), inplace = True)
```

```
df
```

	lread	lwrite	scall	sread	swrite	fork	exec	rchar	wchar	pgout	...	pgscan	atch	pgin	ppgin	pflt	vflt	runqsz	freemem	freeswa
0	1	0	2147	79	68	0.2	0.20	40671.0	53995.0	0.00	...	0.00	0.0	1.60	2.60	16.00	26.40	CPU_Bound	4670	173094
1	0	0	170	18	21	0.2	0.20	448.0	8385.0	0.00	...	0.00	0.0	0.00	0.00	15.63	16.83	Not_CPU_Bound	7278	186900
2	15	3	2162	159	119	2.0	2.40	125473.5	31950.0	0.00	...	0.00	1.2	6.00	9.40	150.20	220.20	Not_CPU_Bound	702	102123
3	0	0	160	12	16	0.2	0.20	125473.5	8670.0	0.00	...	0.00	0.0	0.20	0.20	15.60	16.80	Not_CPU_Bound	7248	186370
4	5	1	330	39	38	0.4	0.40	125473.5	12185.0	0.00	...	0.00	0.0	1.00	1.20	37.80	47.60	Not_CPU_Bound	633	176025
...
8187	16	12	3009	360	244	1.6	5.81	405250.0	85282.0	8.02	...	55.11	0.6	35.87	47.90	139.28	270.74	CPU_Bound	387	98664
8188	4	0	1596	170	146	2.4	1.80	89489.0	41764.0	3.80	...	0.20	0.8	3.80	4.40	122.40	212.60	Not_CPU_Bound	263	105574
8189	16	5	3116	289	190	0.6	0.60	325948.0	52640.0	0.40	...	0.00	0.4	28.40	45.20	60.20	219.80	Not_CPU_Bound	400	96910
8190	32	45	5180	254	179	1.2	1.20	62571.0	29505.0	1.40	...	18.04	0.4	23.05	24.25	93.19	202.81	CPU_Bound	141	102245
8191	2	0	985	55	46	1.6	4.80	111111.0	22256.0	0.00	...	0.00	0.2	3.40	6.20	91.80	110.00	CPU_Bound	659	175651

1192 rows × 22 columns

ii. Check for the Zero values and understand the importance of that data point. Please provide your comments on if we need to change them (impute) or drop them.

```
df.isin([0]).sum()
```

```
lread      675
lwrite     2684
scall       0
sread       0
swrite       0
fork        21
exec        21
rchar       0
wchar       0
pgout      4878
ppgout     4878
pgfree     4869
pgscan     6448
atch       4575
pgin       1220
ppgin      1220
pflt        3
vflt        0
runqsz      0
freemem     0
freeswap    0
usr         283
dtype: int64
```

We can see from the above result we can see how many of the variables have the 0 values.

c. Working with Outliers

i. Check for outliers and provide comments

Check Outliers in the HTML file.

On performing our code we can see that almost every variable has an outlier(refer from HTML file.)

'lread', 'sread', 'swrite', rchar, wchar,exec, atch 'pgout', 'ppgout', 'pgfree', 'pgin', 'ppgin', 'pflt', 'vflt', freemem, freeswap are all right skewed and have the mode is often less than the median, which is less than the mean.

Pg scan has no distribution hence no skewness determined.

Freeswap and usr have a left skewed.

ii. Perform outlier treatment (only if required)

```
def remove_outlier(column):
```

```
    sorted(column)
```

```
    q1=df[column].quantile(0.25)
```

```
    q3=df[column].quantile(0.75)
```

```
    iqr=q3-q1
```

```
    lower=q1-1.5*iqr
```

```
    upper=q3+1.5*iqr
```

```
    return lower,upper
```

```
or i in nums:
```

```
    lower,upper=remove_outlier(i)
```

```
    df[i]=np.where(df[i]>upper,upper,df[i])
```

```
    df[i]=np.where(df[i]<lower,lower,df[i])
```

```
for i in nums:
```

```
    sns.boxplot(df[i],showmeans=True)
```

```
    plt.show()
```

We create a function to remove the outlier and use quartiles to treat the inconsistency if any. From the above code we can treat the outlier.(Can see the corrected outlier treatment in the HTML File)

d. Scaling the data

- i. Only if scaling is necessary, please perform the same and provide your reasoning

```
scaler = MinMaxScaler()
```

```
columns=['lread','lwrite','scall','sread','swrite','fork','exec','rchar','wchar','pgout','ppgout','pgscan','pg  
free','atch','pgin','ppgin','pflt','vflt','freemem','freeswap','usr']
```

```
df_scaled = scaler.fit_transform(df[columns].to_numpy())
```

```
df_scaled = pd.DataFrame(df_scaled, columns=columns)
```

```
print("Scaled Dataset Using MinMaxScaler")
```

```
df_scaled['runqsz']=df['runqsz']
```

```
df_scaled.head().T
```

index	Cb1 Bonuq	io1 Cb1 Bonuq	io1 Cb1 Bonuq	io1 Cb1 Bonuq	io1 Cb1 Bonuq
net	0.883333	0.878881	0.88	0.813333	0.18
leewind	0.110255	0.835388	0.142528	0.858888	0.183881
leewind	1.0	1.0	0.140258	1.0	0.15284
alt	0.048888	0.058833	0.385011	0.058818	0.084483
b11	0.04458	0.043531	0.142481	0.043424	0.104284
b11u	0.011381	0.0	0.518185	0.008885	0.038114
b11u	0.088048	0.0	0.528183	0.008208	0.045281
g11	0.0	0.0	0.8	0.0	0.0
b11ee	0.0	0.0	0.0	0.0	0.0
b12c2u	0.0	0.0	0.0	0.0	0.0
b18on1	0.0	0.0	0.0	0.0	0.0
b18on1	0.0	0.0	0.0	0.0	0.0
w111	0.558111	0.030021	0.135804	0.031304	0.048815
1111	0.088118	0.000518	0.50483	0.50483	0.50483
exec	0.058881	0.058881	0.388508	0.058881	0.088104
1011	0.040818	0.040818	0.408183	0.040818	0.081833
21111	0.188818	0.038184	0.310548	0.054831	0.088813
21821	0.158118	0.051333	0.515	0.010881	0.088813
2c11	0.302152	0.008424	0.301812	0.001824	0.033123
11111	0.0	0.0	0.45	0.0	0.04
11821	0.051511	0.0	0.318148	0.0	0.108383

In this method, we convert variables with different scales of measurements into a single scale. First we check what are the values that needs to be changed. Standard Scaler normalizes the data using the formula $(x - \text{mean}) / \text{standard deviation}$.

e. Working with Categorical variables

i. Identify the categorical data given as part of the data set?

```
df_scaled.runqsz.unique()
```

```
array(['CPU_Bound', 'Not_CPU_Bound'], dtype=object)
```

```
df = pd.get_dummies(df, prefix='CPU_Bound', columns=['runqsz'])
```

```
df.head()
```

lread	lwrite	scall	sread	swrite	fork	exec	rchar	wchar	pgout	...	pgscan	atch	pgin	ppgin	pflt	vflt	runqsz	freemem	freeswap	usi
1	0	2147	79	68	0.2	0.2	40671.0	53995.0	0.0	...	0.0	0.0	1.6	2.6	16.00	26.40	CPU_Bound	4670	1730946	95
0	0	170	18	21	0.2	0.2	448.0	8385.0	0.0	...	0.0	0.0	0.0	0.0	15.63	16.83	Not_CPU_Bound	7278	1869002	97
15	3	2162	159	119	2.0	2.4	125473.5	31950.0	0.0	...	0.0	1.2	6.0	9.4	150.20	220.20	Not_CPU_Bound	702	1021237	87
0	0	160	12	16	0.2	0.2	125473.5	8670.0	0.0	...	0.0	0.0	0.2	0.2	15.60	16.80	Not_CPU_Bound	7248	1863704	96
5	1	330	39	38	0.4	0.4	125473.5	12185.0	0.0	...	0.0	0.0	1.0	1.2	37.80	47.60	Not_CPU_Bound	633	1760253	90

On performing the given code we can see that , runqsz has the categorical values that needs to be changed. Hence perform get dummies function to replace these values with one and 0.(Given in the HTML file)

ii.) Perform encoding and provide detailed comments and reasoning for the encoding approach.

Most of the machine learning models are designed to work on numeric data. Hence, we need to convert categorical text data into numerical data for model building

One-Hot-Encoding is used to create dummy variables to replace the categories in a categorical variable into features of each category and represent it using 1 or 0 based on the presence or absence of the categorical value in the record

3. Perform univariate, bivariate & Multivariate analysis

a. Perform Univariate analysis for each variable and write comments

We have performed univariate analysis for each variable(in the HTML file.) For lread its is positively skewed and greater frequency. For lwrite its similar to lread but its less frequent. For all of the graphs itself we can see an unique peak in the beginning and then a drop which is constant and then a sudden rise as well. For pg scan we can see there is no data to be distributed.

b. Perform bivariate analysis and make necessary inference about the relation between the variables.


Refer html file for the plot.


```
sns.pairplot(df_scaled, diag_kind='kde',size = 3)
plt.show()
```

c. Perform Multivariate analysis and make necessary inferences about the relation between variables.

d. Check Covariance and Correlation and identify positively and negatively correlated variables.

```
plt.figure(figsize=(20,20))
sns.heatmap(df.corr(), annot=True)
```





e. Identify the variables which has multicollinearity. Check for multi collinearity and drop the variables.

```
]: vif_data = pd.DataFrame()
vif_data["feature"] = X.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                    for i in range(len(X.columns))]
cols_to_remove=list(vif_data.loc[vif_data["VIF"]>=5]['feature'])
print(cols_to_remove)

vif_data.sort_values('VIF', ascending=False)
print(vif_data)
```

	feature	VIF
0	lread	5.232535
1	lwrite	4.260737
2	scall	2.989088
3	sread	6.517014
4	swrite	5.639798
5	fork	12.947438
6	exec	3.137525
7	rchar	2.087446
8	wchar	1.605667
9	pgout	11.474544
10	ppgout	30.685098
11	pgfree	17.175122
12	pgscan	NaN
13	atch	1.859658
14	pgin	13.733882
15	ppgin	13.962094
16	pflt	11.536918
17	vflt	15.204685
18	freemem	1.969949
19	freeswap	1.825789
20	CPU_Bound_CPU_Bound	13.213383
21	CPU_Bound_Not_CPU_Bound	13.225624

```
: new_df=X.drop(['lread', 'sread', 'swrite', 'fork', 'pgout', 'ppgout', 'pgfree', 'pgin', 'ppgin', 'pflt', 'vflt', 'CPU_Bound_CPU_Bound', 'CPU_Bound_Not_CPU_Bound'])
new_df
```

```
:
```

	lwrite	scall	exec	rchar	wchar	pgscan	atch	freemem	freeswap
0	0.0	2147.0	0.20	40671.0	53995.0	0.0	0.0	4659.125	1730946.0
1	0.0	170.0	0.20	448.0	8385.0	0.0	0.0	4659.125	1869002.0
2	3.0	2162.0	2.40	125473.5	31950.0	0.0	1.2	702.000	1021237.0
3	0.0	160.0	0.20	125473.5	8670.0	0.0	0.0	4659.125	1863704.0
4	1.0	330.0	0.40	125473.5	12185.0	0.0	0.0	633.000	1760253.0
...
8187	12.0	3009.0	5.81	405250.0	85282.0	0.0	0.6	387.000	986647.0
8188	0.0	1596.0	1.80	89489.0	41764.0	0.0	0.8	263.000	1055742.0
8189	5.0	3116.0	0.60	325948.0	52640.0	0.0	0.4	400.000	969106.0
8190	25.0	5180.0	1.20	62571.0	29505.0	0.0	0.4	141.000	1022458.0
8191	0.0	985.0	4.80	111111.0	22256.0	0.0	0.2	659.000	1756514.0

8192 rows × 9 columns

First we check for multicollinearity and remove all the features that are multicollinear using variance inflation factor. Mathematically, the VIF for a regression model variable is equal to the ratio of the overall model variance to the variance of a model that includes only that single independent variable.

4. Building a Linear Regression Model

a. Prepare data for model building

Preparing the data for model building, we drop the values of `usr` and copy it in another dataset so that we can use this data frame to train and test our model.

```
X = df.drop('usr', axis=1)
```

```
y = df[['usr']]
```

We get the table that we need to test and train.(refer to HTML file.)

b. Build linear regression model.

```
|: X = new_df  
y = df[['usr']]
```

```
|: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25 , random_state=1)
```

```
|: model.predict(X_train)
```

```
|: 5745      68.397776  
1930      53.316204  
5622      60.481046  
5340     102.002734  
2255      98.826406  
...  
7935      74.811305  
5192      96.452824  
3980      69.561082  
235       70.027195  
5157      94.588283  
Length: 6144, dtype: float64
```

```
|: regression_model = LinearRegression()  
regression_model.fit(X_train, y_train)
```

Following all the required codes to achieve the model to further find the required measures to check accuracy of our model.

c. Find the features that add value to the model. Identify the list of Variables which highly impact the prediction based on the correlation Metrix given for regression (target variable)

```

: vif = [variance_inflation_factor(X.values, ix) for ix in range(X.shape[1])]

: i=0
  for column in X.columns:
    if i < 15:
      print (column, "---->", vif[i])
      i = i+1

lread ----> 5.232534882470838
lwrite ----> 4.260737491684919
scall ----> 2.9890877339622937
sread ----> 6.517014292159949
swrite ----> 5.639797517443181
fork ----> 12.947438478657318
exec ----> 3.1375249704246078
rchar ----> 2.087445737239059
wchar ----> 1.6056667145387191
pgout ----> 11.474543882811679
ppgout ----> 30.685097541787272
pgfree ----> 17.17512180642015
pgscan ----> nan
atch ----> 1.8596581052295806
pgin ----> 13.733882306297474

```

Linear Regression Model

Again we can find the best correlated values that suit for our model. First we check for multicollinearity and remove all the features that are multicollinear using variance inflation factor. Mathematically, the VIF for a regression model variable is equal to the ratio of the overall model variance to the variance of a model that includes only that single independent variable.


d. Print test and train results with all variables and best fit line.

```

sns.set(rc = {'figure.figsize':(7,7)})
ax=sns.regplot(x=y_test,y=y_pred,ci=None,color = 'blue');
ax.set_xlabel('Actual')
ax.set_ylabel('Predicted')

```

Text(0, 0.5, 'Predicted')



Heavily clustered but in a consistent manner, but the growth seems to be very consistent.

5. Model Performance :

a. Check for the performance measures for linear regression (Hint: RMSE, R square, etc.)

```
: # Let us check the intercept for the model
intercept = regression_model.intercept_[0]

print("The intercept for our model is {}".format(intercept))

The intercept for our model is 52.703408616461346

: # R square on training data
regression_model.score(X_train, y_train)

: 0.5513661898690928

: #R square on testing data
regression_model.score(X_test, y_test)

: 0.5447312405803018

: #RMSE on training data
predicted_train=regression_model.fit(X_train, y_train).predict(X_train)
np.sqrt(mean_squared_error(y_train,predicted_train))

: 12.16989876561295

: ##RMSE on testing data
predicted_test=regression_model.fit(X_train, y_train).predict(X_test)
np.sqrt(mean_squared_error(y_test,predicted_test))

: 12.870554763335232

: # concatenate X and y into a single dataframe
data_train = pd.concat([X_train, y_train], axis=1)
data_test=pd.concat([X_test,y_test],axis=1)
data_train.head()
```

```
:
      lwrite  scall  exec    rchar    wchar  pgscan  atch  freemem  freeswap  usr
5745   10.0  2007.0   0.60  32665.0  49643.0     0.0   0.6    329.0   989029.0   94
1930    4.0   837.0   0.80   6255.0  23670.0     0.0   0.0   3052.0  1013758.0   96
5622    2.0  2227.0   6.70 108370.0  28568.0     0.0   0.2    314.0  1108418.0   77
5340    0.0  2132.0   0.20   90813.0  28590.0     0.0   0.0    488.0  1742493.0   93
2255    0.0  3517.0   2.99  310439.0  214462.0     0.0   0.2    420.0  1547657.0   86
```

From above relations we can see that my RMSE value is 12.16 for test and train almost. For R square the value is 0.544.

b. Experiment with data transformation and suggest if we can improve the model performance.

```
100 - (rmse/y_train.mean())*100

usr      85.523778
dtype: float64
```

To improve the performance of this model we can decrease the RMSE values and increase the R² mean value.